

Name	What did you do last week?	What do you plan to do this week?	What do you plan to do this week?
Aryaman Jalali	Revised the abstract and introduction based on the feedback, set up the shared GitHub repository with folders for ETL, anomaly detection, and dashboard code, and documented the initial project structure for the team.	Design and finalize the anomaly output schema (fields like anomaly type, confidence, baseline stats, etc.), set up the storage layer for the anomalies table, and build a small prototype dashboard/notebook view that can list and filter anomalies for a given month.	No blockers: mainly syncing with the team on final field names and anomaly-type labels before wiring them into the code.
Harshith Keshavamurthy	Added brief descriptions and value ranges for all input and output features in the proposal, created an exploratory Spark notebook, and computed basic statistics on a sample month of the Wikipedia clickstream data.	Implement baseline statistics for each edge using Spark (median, MAD, and rolling mean/standard deviation), start coding the traffic-spike detection logic using robust Z-scores, and draft a short note on the evaluation metrics we plan to use (e.g., Precision@K).	No blockers: currently focusing on choosing sensible thresholds and sampling strategies so the baseline jobs stay efficient.
Dirgha Pareshkumar Jivani	Refined the success metrics into more fine-grained, quantitative goals, configured the Spark environment, and verified that one month of clickstream data can be ingested and processed end-to-end as an initial ETL test.	Implement the full ETL pipeline for multiple months of clickstream data, including reading TSVs, cleaning and filtering edges, normalizing page identifiers, and writing the cleaned dataset to Parquet, and document the ETL steps for the rest of the team.	No blockers: just monitoring job runtimes and storage usage as more months are added to ensure the pipeline scales smoothly.